Simultaneous Elicitation of Scoring Rule and Agent Preferences for Robust Winner Determination

B. Napolitano[1], O. Cailloux[1] and P. Viappiani[2]

[1] Université Paris-Dauphine, Université PSL, CNRS, LAMSADE
[2] LIP6, UMR 7606, CNRS and Sorbonne Université

ADT 2021 - 04 November 2021

Introducing the problem

Agents = { , ,  }, Alternatives = { , ,  }, Chair =  $\Rightarrow$ Positional Scoring Rule

Introducing the problem

Agents $= \{\ ,\ ,\ \}$, Alternatives $= \{\ ,\ ,\ \}$, Chair $=$ $\Rightarrow$ Positional Scoring Rule

Introducing the problem

Agents $= \{$, , $\}$, Alternatives $= \{$, , $\}$, Chair $=$ $\Rightarrow$ Positional Scoring Rule

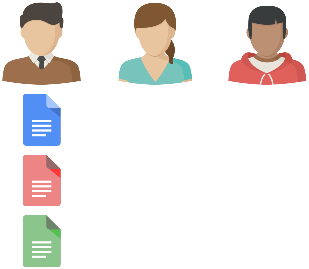

Introducing the problem

Agents = { , , }, Alternatives = { , , }, Chair = $\Rightarrow$ Positional Scoring Rule

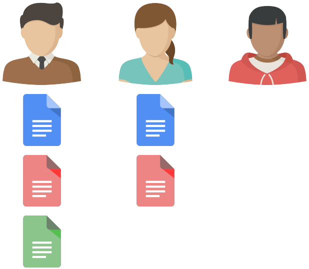

Introducing the problem

Agents $= \{$, , $\}$, Alternatives $= \{$, , $\}$, Chair $=$ $\Rightarrow$ Positional Scoring Rule

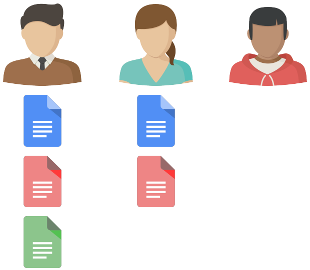

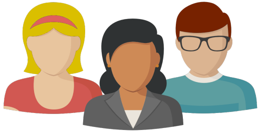

Introducing the problem

Agents $= \{\ ,\ ,\ \}$, Alternatives $= \{\ ,\ ,\ \}$, Chair $=$ $\Rightarrow$ Positional Scoring Rule

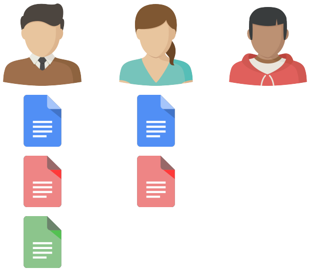

$\text{weight}(1^{\text{st}}) \geq 2 \cdot \text{weight}(2^{\text{nd}})$

Introducing the problem

Setting: Incompletely specified preferences and social choice rule

Introducing the problem

Setting: Incompletely specified preferences and social choice rule

- Agents: difficult or costly to order *all* alternatives
- Chair: difficult to *specify* a voting rule precisely

Introducing the problem

Setting: Incompletely specified preferences and social choice rule

Goal: Reduce uncertainty, inferring (*eliciting*) incrementally and simultaneously the true preferences of agents and chair to quickly converge to an optimal or a near-optimal alternative

Introducing the problem

Setting: Incompletely specified preferences and social choice rule

Goal: Reduce uncertainty, inferring (*eliciting*) incrementally and simultaneously the true preferences of agents and chair to quickly converge to an optimal or a near-optimal alternative

Approach:

- Develop query strategies that interleave questions to the chair and to the agents
- Use *Minimax regret* to measure the quality of those strategies

Related Works

Incomplete profile

- and known weights: Minimax regret to produce a robust winner approximation (*Lu and Boutilier 2011*, [2]; *Boutilier et al. 2006*, [1])

Uncertain weights

- and complete profile: dominance relations derived to eliminate alternatives always less preferred than others (*Stein et al. 1994*, [3])
- in positional scoring rules (*Viappiani 2018*, [4])

Assumptions

- We consider *Positional Scoring Rules*, which attach weights to positions according to a scoring vector W
- We assume W to be *convex*

$$W_r - W_{r+1} \geq W_{r+1} - W_{r+2}$$

for all positions r, and that $W_1 = 1$ and $W_m = 0$

Notation

A	alternatives, $\lvert A \rvert = m$
N	agents (*voters*)
$P = (\succ_j,\ j \in N),\ P \in \mathcal{P}$	complete preferences profile
$W = (W_r,\ 1 \leq r \leq m),\ W \in \mathcal{W}$	(convex) scoring vector that the chair has in mind

Notation

$$\begin{aligned} A &\ \text{alternatives, } |A| = m \\ N &\ \text{agents (\textit{voters})} \\ P = (\succ_j,\ j \in N),\ P \in \mathcal{P} &\ \text{complete preferences profile} \\ W = (W_r,\ 1 \le r \le m),\ W \in \mathcal{W} &\ \text{(convex) scoring vector that the chair has in mind} \end{aligned}$$

W defines a Positional Scoring Rule $f_W(P) \subseteq A$ using scores $s^{W,P}(a),\ \forall\ a \in A$

Notation

A alternatives, $|A| = m$

N agents (*voters*)

$P = (\succ_j,\ j \in N),\ P \in \mathcal{P}$ complete preferences profile

$W = (W_r,\ 1 \leq r \leq m),\ W \in \mathcal{W}$ (convex) scoring vector that the chair has in mind

W defines a Positional Scoring Rule $f_W(P) \subseteq A$ using scores $s^{W,P}(a),\ \forall\ a \in A$

P and W exist in the minds of agents and chair but unknown to us

Questions

Two types of questions:

Questions

Two types of questions:

Questions to the agents

Comparison queries that ask a particular agent j to compare two alternatives $a, b \in \mathcal{A}$

$$a \succ_j b \quad ?$$

Questions

Two types of questions:

Questions to the agents

Comparison queries that ask a particular agent j to compare two alternatives $a, b \in \mathcal{A}$

$$a \succ_j b \quad ?$$

Questions to the chair

Queries relating the difference between the importance of consecutive ranks from r to $r + 2$

$$W_r - W_{r+1} \geq \lambda(W_{r+1} - W_{r+2}) \quad ?$$

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

- $C_P \subseteq \mathcal{P}$ constraints on the profile given by the agents

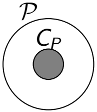

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

- $C_P \subseteq \mathcal{P}$ constraints on the profile given by the agents
- $C_W \subseteq \mathcal{W}$ constraints on the voting rule given by the chair

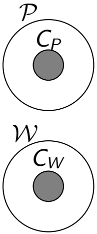

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

- $C_P \subseteq \mathcal{P}$ constraints on the profile given by the agents
- $C_W \subseteq \mathcal{W}$ constraints on the voting rule given by the chair

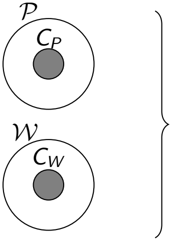

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

- $C_P \subseteq \mathcal{P}$ constraints on the profile given by the agents
- $C_W \subseteq \mathcal{W}$ constraints on the voting rule given by the chair

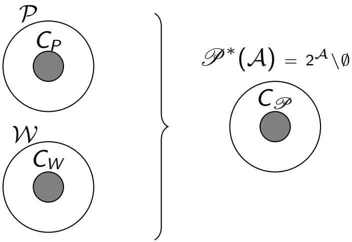

Current Knowledge

The answers to these questions define C_P and C_W that is our knowledge about P and W

- $C_P \subseteq \mathcal{P}$ constraints on the profile given by the agents
- $C_W \subseteq \mathcal{W}$ constraints on the voting rule given by the chair

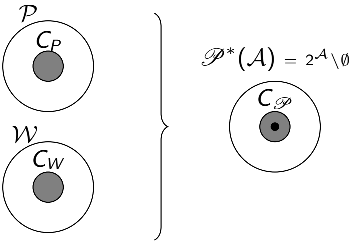

Minimax Regret

Given $C_P \subseteq \mathcal{P}$ and $C_W \subseteq \mathcal{W}$:

$$\mathrm{PMR}^{C_P, C_W}(a, b) = \max_{P \in C_P, W \in C_W} s^{P,W}(b) - s^{P,W}(a)$$

is the maximum difference of score between a and b under all possible realizations of the full profile *and* weights

Minimax Regret

Given $C_P \subseteq \mathcal{P}$ and $C_W \subseteq \mathcal{W}$:

$$\mathrm{PMR}^{C_P,C_W}(a,b) = \max_{P \in C_P, W \in C_W} s^{P,W}(b) - s^{P,W}(a)$$

is the maximum difference of score between a and b under all possible realizations of the full profile *and* weights

We care about the worst case loss: *maximal regret* between a chosen alternative a and best real alternative b

$$\mathrm{MR}^{C_P,C_W}(a) = \max_{b \in \mathcal{A}} \mathrm{PMR}^{C_P,C_W}(a,b)$$

Minimax Regret

Given $C_P \subseteq \mathcal{P}$ and $C_W \subseteq \mathcal{W}$:

$$\mathrm{PMR}^{C_P,C_W}(a,b) = \max_{P\in C_P, W\in C_W} s^{P,W}(b) - s^{P,W}(a)$$

is the maximum difference of score between a and b under all possible realizations of the full profile *and* weights

We care about the worst case loss: *maximal regret* between a chosen alternative a and best real alternative b

$$\mathrm{MR}^{C_P,C_W}(a) = \max_{b\in\mathcal{A}} \mathrm{PMR}^{C_P,C_W}(a,b)$$

We select the alternative which *minimizes* the maximal regret

$$\mathrm{MMR}^{C_P,C_W} = \min_{a\in\mathcal{A}} \mathrm{MR}^{C_P,C_W}(a)$$

Pairwise Max Regret Computation

The computation of $\mathrm{PMR}^{C_P, C_W}($, $)$ can be seen as a game in which an adversary both:

Pairwise Max Regret Computation

The computation of $\mathrm{PMR}^{C_P,C_W}(\ ,\ )$ can be seen as a game in which an adversary both:

- **chooses a complete profile P** $\in \mathcal{P}$

Pairwise Max Regret Computation

The computation of $\mathrm{PMR}^{C_P,C_W}($, $)$ can be seen as a game in which an adversary both:

- **chooses a complete profile P** $\in \mathcal{P}$

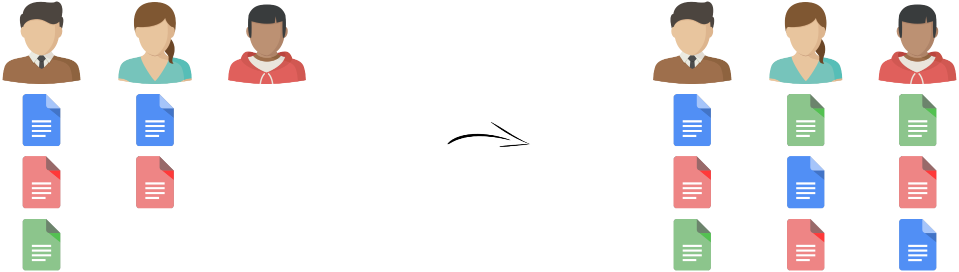

- **chooses a feasible weight vector W** $\in \mathcal{W}$

$$(\mathbf{1}, ?, \mathbf{0}) \longrightarrow (\mathbf{1}, \mathbf{0}, \mathbf{0})$$

in order to maximize the difference of scores

Elicitation strategies

At each step, the strategy selects a question to ask either to one of the agents about her preferences or to the chair about the voting rule

Elicitation strategies

At each step, the strategy selects a question to ask either to one of the agents about her preferences or to the chair about the voting rule

The termination condition could be:

Elicitation strategies

At each step, the strategy selects a question to ask either to one of the agents about her preferences or to the chair about the voting rule

The termination condition could be:

- when the minimax regret is lower than a threshold

Elicitation strategies

At each step, the strategy selects a question to ask either to one of the agents about her preferences or to the chair about the voting rule

The termination condition could be:

- when the minimax regret is lower than a threshold
- when the minimax regret is zero

Assume that a question leads to the possible new knowledge states (C_P^1, C_W^1) and (C_P^2, C_W^2) depending on the answer, then the badness of the question in the worst case is:

$$\max_{i=1,2} \mathrm{MMR}(C_P^i, C_W^i)$$

Elicitation strategies

Pessimistic Strategy

Assume that a question leads to the possible new knowledge states (C_P^1, C_W^1) and (C_P^2, C_W^2) depending on the answer, then the badness of the question in the worst case is:

$$\max_{i=1,2} \mathrm{MMR}(C_P^i, C_W^i)$$

The pessimistic strategy selects the question that leads to minimal regret in the worst case from a set of $n+1$ candidate questions

Assume that a question leads to the possible new knowledge states (C_P^1, C_W^1) and (C_P^2, C_W^2) depending on the answer, then the badness of the question in the worst case is:

$$\max_{i=1,2} \mathrm{MMR}(C_P^i, C_W^i)$$

The pessimistic strategy selects the question that leads to minimal regret in the worst case from a set of $n+1$ candidate questions

Note:

if the maximal MMR of two questions are equal, then prefers the one with the lowest MMR values associated to the opposite answer

Let $(a^*, \bar{b}, \bar{P}, \bar{W})$ be the current solution of the minimax regret

We select $n+1$ candidate questions:

- **One question per agent:** For each agent i, either:
  - $a^* \succ_j^{\bar{P}} \bar{b}$: we ask about an incomparable alternative that can be placed above a^* by the adversary to increase PMR(a^*,$\bar{b}$)
  - $\bar{b} \succ_j^{\bar{P}} a^*$: we ask about an incomparable alternative that can be placed between a^* and $\bar{b}$ by the adversary to increase PMR(a^*,$\bar{b}$)
  - a^* and $\bar{b}$ are incomparable: we ask to compare them
- **One question to the chair:** Consider W_τ the weight vector that minimize the PMR in the worst case.

  We ask about the position $r = \underset{i=[\![1,m-1]\!]}{\arg\max} |\bar{W}(i) - W_\tau(i)|$

Empirical Evaluation

Pessimistic for different datasets

Figure: Average MMR (normalized by *n*) after *k* questions with Pessimistic strategy for different datasets.

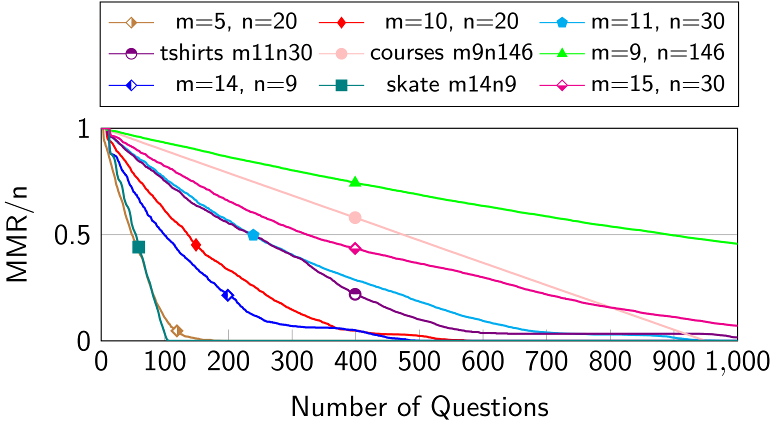

Table: Questions asked by Pessimistic strategy on several datasets to reach $\frac{n}{10}$ regret, columns 4 and 5, and zero regret, last two columns.

dataset	m	n	$q_c^{MMR \le n/10}$	$q_a^{MMR \le n/10}$	$q_c^{MMR=0}$	$q_a^{MMR=0}$
m5n20	5	20	0.0	[4.3 | 5.0 | 5.8]	5.3	[5.4 | 6.2 | 7.2]
m10n20	10	20	0.0	[13.9 | 16.1 | 18.4]	32.0	[19.7 | 21.8 | 24.7]
m11n30	11	30	0.0	[16.6 | 19.0 | 22.3]	45.2	[23.1 | 25.7 | 28.9]
tshirts	11	30	0.0	[13.1 | 16.6 | 19.6]	43.2	[28.2 | 32.0 | 35.6]
courses	9	146	0.0	[6.0 | 7.0 | 7.0]	0.0	[6.8 | 7.0 | 7.0]
m14n9	14	9	5.4	[30.3 | 33.5 | 36.7]	64.1	[37.6 | 40.5 | 44.3]
skate	14	9	0.0	[11.4 | 11.6 | 12.3]	0.0	[11.5 | 11.8 | 12.8]
m15n30	15	30	0.0	[25.0 | 29.5 | 33.7]		

Empirical Evaluation

Pessimistic chair first and then agents (and vice-versa)

Table: Average MMR in problems of size $(10, 20)$ after 500 questions, among which q_c to the chair.

q_c	2 ph. ca $\pm$ sd	2 ph. ac $\pm$ sd
0	0.6 ± 0.5	0.6 ± 0.5
15	0.5 ± 0.5	0.5 ± 0.5
30	0.3 ± 0.5	0.3 ± 0.4
50	0.0 ± 0.1	0.0 ± 0.1
100	0.1 ± 0.2	0.1 ± 0.1
200	2.3 ± 1.4	2.1 ± 1.8
300	5.2 ± 2.4	6.8 ± 0.6
400	10.9 ± 0.9	12.2 ± 1.0
500	20.0 ± 0.0	20.0 ± 0.0

Thank You!

Craig Boutilier, Relu Patrascu, Pascal Poupart, and Dale Schuurmans.
Constraint-based optimization and utility elicitation using the minimax decision criterion.
Artificial Intelligence, 2006.

Tyler Lu and Craig Boutilier.
Robust approximation and incremental elicitation in voting protocols.
In *Proc. of IJCAI'11*, 2011.

William E. Stein, Philip J. Mizzi, and Roger C. Pfaffenberger.
A stochastic dominance analysis of ranked voting systems with scoring.
EJOR, 1994.

Paolo Viappiani.
Positional scoring rules with uncertain weights.
In *Scalable Uncertainty Management*, 2018.